

Temporal fusion

candidate



motion, entropy, attention
stable patches keep $t-1$ tokens
capped, keyframes bound drift

Depth pruning

candidate



Block Influence ranking
ends protected, none adjacent
layers removed, not bypassed

Guarded reuse


candidate



all pass: reuse a_{t-1}

any fails: dense call

image and action gates before the call
pass: skip the call, reuse a_{t-1}



Foveation

control



sharp disc, blurred
periphery, same
token count

Protocol

dense reference under fused attention, weights frozen
matched episodes paired by seed, on every backbone
and environment with a checkpoint
positive only if latency falls, own signal cost included,
or success rises, the other within a margin

Action repeat

control



k steps

hold each action k steps, no gate