


1) When the policy runs

model calls per environment step

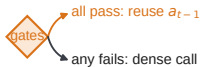
Action repeathold each action k steps, no gate**2) What it is shown**

pixels or visual tokens

Fixed foveationsharp disc of keep ratio ρ ,
blur outside, token count same**3) How much of the decoder runs**

decoder layers per call

no control on this axis

Guarded reuseskip the call only when the image is stable,
the last two dense actions agree and still
command motion, capped**Temporal fusion**stable tokens carried from $t-1$,
capped, entropy and attention
protect, keyframes bound drift**Depth pruning**Block Influence ranking, protected
regions, no adjacent removals,
set fixed on a calibration split**Protocol****Reference**released policy,
fused attention**Backbones**

six, 0.5B to 7B

EnvironmentsWidowX Bridge,
Google Robot,
LIBERO suites**Episodes**matched seeds,
paired test**Positive gate**latency down or
success up, other
within a margin