

1) When the policy runs

model calls per environment step

Action repeat

k steps
hold each action k steps, no gate

2) What it is shown

pixels or visual tokens


Foveation

sharp disc of keep ratio ρ ,
blur outside, token count same

3) How much of the decoder runs

decoder layers per call

no control on this axis

Guarded reuse

before the call: skip it and emit a_{t-1}
only when the image is stable, the last two
dense actions agree and command motion, capped

Temporal fusion

patches low in motion, entropy
and optional attention take the
 $t-1$ token, ring, cap, keyframes

Depth pruning

lowest Block Influence layers
removed, ends kept, none adjacent,
set on a calibration split

Protocol**Reference**

released policy,
fused attention

Grid

every backbone on
each environment
with a checkpoint

Environments

WidowX Bridge,
Google Robot,
LIBERO suites

Episodes

matched seeds,
paired by seed

Positive gate

latency down, own
signal cost in, or
success up, other
within a margin