

1) When the policy runs

model calls per environment step


Action repeat k stepshold each action k steps, no gate**2) What it is shown**

pixels or visual tokens

Foveationsharp disc of keep ratio ρ ,
blur outside, token count same**3) How much of the decoder runs**

decoder layers per call

no control on this axis

Guarded reuse

before the call: skip it and emit a_{t-1}
only when the image is stable, the last two
dense actions agree and command motion, capped

Temporal fusionpatches low in motion, entropy
and optional attention take the
 $t-1$ token, ring, cap, keyframes**Depth pruning**lowest Block Influence layers
removed, ends kept, none adjacent,
set on a calibration split**Protocol****Reference**released policy,
fused attention**Grid**every backbone on
each environment
with a checkpoint**Environments**WidowX Bridge,
Google Robot,
LIBERO suites**Episodes**matched seeds,
paired by seed**Positive gate**latency down, own
signal cost in, or
success up, other
within a margin