


**1) When the policy runs**

model calls per environment step

**Action repeat** $k$  stepshold each action  $k$  steps, no gate**Guarded reuse**

before the call: skip it and emit  $a_{t-1}$   
 only when the image is stable, the last two dense actions agree and command motion, capped

**2) What it is shown**

pixels or visual tokens

**Foveation**

sharp disc of keep ratio  $\rho$ ,  
 blur outside, token count same

**Temporal fusion**

patches low in motion, entropy  
 and optional attention take the  
 $t-1$  token, ring, cap, keyframes

**3) How much of the decoder runs**

decoder layers per call

no control on this axis

**Depth pruning**

lowest Block Influence layers  
 removed, ends kept, none adjacent,  
 set on a calibration split

**Protocol****Reference**

released policy,  
 fused attention

**Grid**

every backbone on  
 each environment  
 with a checkpoint

**Environments**

WidowX Bridge,  
 Google Robot,  
 LIBERO suites

**Episodes**

matched seeds,  
 paired by seed

**Positive gate**

latency down, own  
 signal cost in, or  
 success up, other  
 within a margin