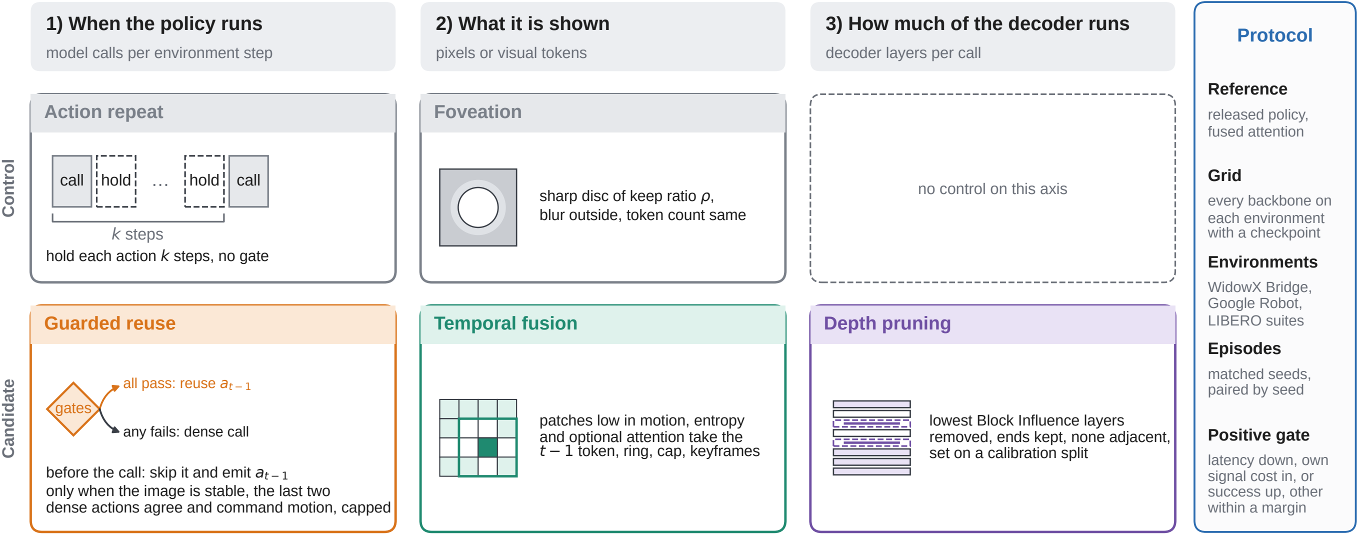
1) When the policy runs
model calls per environment step
2) What it is shown
pixels or visual tokens
3) How much of the decoder runs
decoder layers per call
Protocol
Control
Action repeat
call
hold
…
hold
call
k steps
hold each action k steps, no gate
Foveation
sharp disc of keep ratio ρ,
blur outside, token count same
no control on this axis
Reference
released policy,
fused attention
Grid
every backbone on
each environment
with a checkpoint
Environments
WidowX Bridge,
Google Robot,
LIBERO suites
Episodes
matched seeds,
paired by seed
Positive gate
latency down, own
signal cost in, or
success up, other
within a margin
Candidate
Guarded reuse
gates
all pass: reuse a_{t-1}
any fails: dense call
before the call: skip it and emit a_{t-1}
only when the image is stable, the last two
dense actions agree and command motion, capped
Temporal fusion
patches low in motion, entropy
and optional attention take the
$t-1$ token, ring, cap, keyframes
Depth pruning
lowest Block Influence layers
removed, ends kept, none adjacent,
set on a calibration split