



1) When the policy runs

model calls per environment step

Action repeat k stepshold each action k steps, no gate**Guarded reuse**

before the call: skip it and emit a_{t-1}
 only when the image is stable, the last two dense actions agree and command motion, capped

2) What it is shown

pixels or visual tokens

Foveation


sharp disc of keep ratio ρ ,
 blur outside, token count same

Temporal fusion

patches low in motion, entropy
 and optional attention take the
 $t-1$ token, ring, cap, keyframes

3) How much of the decoder runs

decoder layers per call

**Depth pruning**

lowest Block Influence layers
 removed, ends kept, none adjacent,
 set on a calibration split

Protocol**Reference**released policy,
fused attention**Grid**every backbone on
each environment
with a checkpoint**Environments**WidowX Bridge,
Google Robot,
LIBERO suites**Episodes**matched seeds,
paired by seed**Positive gate**latency down, own
signal cost in, or
success up, other
within a margin