

**1) When the policy runs**

model calls per environment step


**Action repeat** $k$  stepshold each action  $k$  steps, no gate**2) What it is shown**

pixels or visual tokens

**Foveation**sharp disc of keep ratio  $\rho$ ,  
blur outside, token count same**3) How much of the decoder runs**

decoder layers per call

no control on this axis

**Guarded reuse**

before the call: skip it and emit  $a_{t-1}$   
only when the image is stable, the last two  
dense actions agree and command motion, capped

**Temporal fusion**patches low in motion, entropy  
and optional attention take the  
 $t-1$  token, ring, cap, keyframes**Depth pruning**lowest Block Influence layers  
removed, ends kept, none adjacent,  
set on a calibration split**Protocol****Reference**released policy,  
fused attention**Grid**every backbone on  
each environment  
with a checkpoint**Environments**WidowX Bridge,  
Google Robot,  
LIBERO suites**Episodes**matched seeds,  
paired by seed**Positive gate**latency down, own  
signal cost in, or  
success up, other  
within a margin